

June 25, 2026 • Brandon__India • brandon@majikl.com


majiklTrust • MARKET INTELLIGENCE RELEASE • 1.9.35

Market Intelligence The AI Content Agent

Live research in. Governed, on-brand posts out, published to your company page, with a human in command.

alpha.majikl.com

Investor Overview



WHAT IT IS

From a research feed to a published post, autonomously, with a human in command.



Isolated per organization

Multi-tenant SaaS: every tenant gets a private, row-level-secured workspace and its own data.



Grounded in live sources

The agent reads your RSS/Atom feeds and drafts posts from current, corroborated research.



A human approval gate

Nothing reaches the company page until an owner approves it. The gate is server-enforced.



Publishes on schedule

Approved posts publish automatically to your LinkedIn organization page at the right time.



Two roles, one governed workspace.



Dana, Owner
Marketing Director

Runs the workspace. Invites teammates, sets roles, owns the global topics, and is the one who approves what ships.

Manage users

All topics

Approve / reject



Marcus, Member
Content Strategist

Turns topics into posts. Builds angles, generates drafts from research, refines with AI, and submits for approval.

Own topics

Generate & refine

Submit for approval



Priya, Member
Research Analyst

Keeps the sourcing honest. Discovers and validates feeds, tags domains, and tunes the research instructions.


Feed discovery

Research tuning

Sourced accuracy

The content pipeline

Five user-visible states. The approval gate is the control point, and a client can never skip it.



Internal states handled by the publisher, not the client: publishing → posted / failed / blocked.

Bring the team in, and control exactly who can do what.



1

Invite by email

The owner sends an invite from the admin panel; pending invites are tracked and revocable.



2

They join the tenant

The teammate accepts and becomes a member of that organization, and only that one.



3

Assign a role

Grant Owner or Member. Roles map to real permissions: who can manage topics, and who approves.



4

Change or revoke anytime

Promote, demote, or remove access in one click. Synthetic dev-bypass users are blocked outright.



Owner-gated by the `manage_users` permission, blocked for dev-bypass identities, and isolated per tenant by row-level security: three independent gates on one screen.

Define what to talk about, and how to keep it fresh.

1 Create a topic
Name it, describe it, add hashtags, and set scope: a shared global topic or your own personal one.

2 Define content angles
Give the topic a set of angles: the distinct takes the agent can write from.

3 Rotate or pin
The agent auto-rotates angles to avoid repeating itself, or you pin a specific angle for a post.

4 Let AI draft a starter
Not sure where to begin? AI proposes a ready-to-edit topic to get you moving.

Angles are a closed set: arbitrary free text is rejected (fail-closed), so user input never flows into the prompt.

CONTENT ANGLES

- ☒ Contrarian take
- ☐ Data-backed trend
- ☐ Practitioner how-to
- ☐ Myth vs. reality

Point the agent at the right sources, safely.

- 

1

Discover feeds with AI
Describe a topic and the agent suggests relevant RSS/Atom feeds to pull from.
- 

2


Validated before added
Every candidate URL is RSS-validated and SSRF-checked, so only safe, real feeds get in.
- 

3

Tag & instruct
Tag feeds by domain and set research instructions that steer what the agent looks for.
- 

4

Continuous ingestion
The research engine keeps pulling and corroborating articles in the background.



Turn the latest research into a draft in one click.

MR

Marcus
Member

- 

1

Pick topic & angle

Choose what to post about and which angle to write from, or let it auto-select.
- 

2

Draft from research

The agent composes a LinkedIn post grounded in the most recent corroborated articles.
- 

3

Verified metrics, sourced

Numbers are woven in with their source attached. An unsourced figure never ships.
- 

4

Your model, your key

Pick the model that fits; each tenant brings its own AI key (BYOK).



Shape the draft until it sounds like you.



1

Ask AI to rewrite

Tighten it, reframe the hook, or shift the tone. The agent rewrites on request.



2

Edit directly

Prefer the keyboard? Edit the post text inline and keep full control of the wording.



3

Iterate freely

Generate, rewrite, tweak. Loop as many times as it takes to get it right.



4

Every call is traceable

AI requests are logged end-to-end, so what the model did is always auditable.



Refine and Generate share the same disciplined pipeline: the exact request sent to the model is the exact request logged, so the trace can never lie about what shipped.

Give every post the right picture, fetched safely.



1 Pull the lead image

The agent extracts the lead image straight from the source article.




2 Smart OG fallback

No lead image? It falls back to the article's Open-Graph preview image.



3 Fetched through a guarded proxy

Images load via an SSRF-guarded proxy: no server-side request forgery.





4 Preview before it ships

See the post exactly as it will appear on the company page.



Save it, submit it, and route it for sign-off.



Queue it once: the agent posts it for you.

-  **1 Queue for a time**
Put an approved post on the calendar for the date and time you want it live.
-  **2 Track schedule status**
See what's queued, what's gone out, and what needs attention at a glance.
-  **3 Auto-publish to your page**
The scheduler posts to your LinkedIn organization page: no manual step.
-  **4 End-to-end status**
Posted, failed, and blocked are all tracked, so nothing silently disappears.

THE QUEUE

 Scheduled · Tue 9:00	SCHEDULED
 Posted · Mon 9:00	POSTED
 Pending approval	PENDING
 Failed · retry queued	FAILED



NEW

WORKFLOW · REWRITE WITH AI



Marcus
Member

Reshape a draft without starting over.



1 One click, no prompt

Hit Rewrite with AI and it reuses the post's own topic, angle, and the research it was built on.




2 Polishes, never fabricates

The refine pass tightens and reframes the existing copy. It is metric-free, so it cannot invent a new number.



3 Returns to draft for sign-off

Rewriting a post that was awaiting approval drops it back to draft, so a changed post is always re-approved.





4 Gated and traceable

Guarded by the edit_post permission, and every rewrite is logged exactly like a generation.



Rewrite runs the same audited pipeline as Generate, on a separate metric-free prompt, so a polished draft stays as accountable as the original.





NEW

SCHEDULE · PUBLISHING ON A FUTURE DATE



Dana
Owner

Set it today. It posts itself next Tuesday.



1 Start from a ready post

Pick a post that has cleared approval and is ready to go out.



2 Pick a future date & time

Choose Scheduled and set when it should publish. The system rejects any time that isn't in the future.



3 It joins the queue

The post shows a Scheduled status with your chosen time and stays editable right up until it fires.



4 It posts itself

At that moment the scheduler publishes to your LinkedIn org page: posted, or failed / blocked, all tracked.



An optional spacing rule keeps scheduled posts from bunching, separate from the always-on engine that paces autonomous posts (roughly 4 per 10 days, 72h apart).

June 2026

9:00 AM

S	M	T	W	T	F	S
	1	2	3	4	5	6
7	8	9	10	11	12	13
14	15	16	17	18	19	20
21	22	23	24	25	26	27
28	29	30				

☒ Scheduled: Jun 30, 9:00 AM

☐ Today: Jun 25



NEW

WORKFLOW · RESEARCH INSTRUCTIONS



Priya
Member

Tell the agent what counts as a good source.

Research instructions are a few search templates on each topic. They steer what the agent looks for, instead of letting it guess from the topic name.

LEFT EMPTY

The agent derives generic queries from the topic name:

```
AI in cybersecurity  
AI in cybersecurity trends
```

You get: *broad explainers and vendor hype: a vague post with no hard evidence.*

WITH RESEARCH INSTRUCTIONS

You give it templates that target citable material:

```
{{TOPIC_NAME}} breach incident report {{YEAR_RANGE}}  
AI security {{KEYWORDS}} enforcement action  
{{ANGLE}} case study results survey
```


You get: *incident reports, enforcement actions, and studies with numbers: a specific, sourced post.*



Up to five templates, four placeholders: {{TOPIC_NAME}} {{ANGLE}} {{KEYWORDS}} {{YEAR_RANGE}}. The AI drafts them; you review and save. Sharper sourcing means stronger metrics and fewer rejected drafts.

A governed loop that runs continuously

Research never stops, so neither does the pipeline, with human approval as the one deliberate pause.



Hands-free throughput

Posts keep flowing without anyone babysitting the tool.



Always current

Every post reflects the newest corroborated research, not last month's.



One human checkpoint

The owner's approval is the only gate: deliberate, not bureaucratic.



Nothing slips

Failed or blocked posts surface immediately with their status.

Security is the product, not a feature

Zero-Trust from the database to the UI: the moat for selling AI publishing to serious organizations.



Row-level multi-tenancy

Postgres RLS isolates every tenant's data, enforced in the database, not just the app.



Per-tenant BYOK keys

Each organization brings its own AI key; keys are encrypted and never logged.



Encrypted prompt vault

Prompts and secrets sealed with AES-256-GCM + HKDF at rest.



Human approval gate

No post reaches a public page without an owner's sign-off, server-enforced.



SSRF-guarded fetches

Feed discovery and image fetches are validated against server-side request forgery.



Full AI traceability

Every model request is logged exactly as sent, so output is always auditable.

June 25, 2026, • Brandon _ India • brandon@majikl.com

THE OPPORTUNITY

Market intelligence that posts itself, safely.

A governed AI agent that researches, drafts, and publishes on-brand LinkedIn content for every organization on the platform, with the owner always in command.

Research

live, corroborated sources in

alpha.majikl.com

Governance

human approval before publish

Scale

multi-tenant, isolated by design